

# Microservices Migration Patterns

**Armin Balalaie**

Department of Computer Engineering, Sharif University of Technology, Tehran, Iran  
The corresponding author, [armin.balalaie@gmail.com](mailto:armin.balalaie@gmail.com)

**Abbas Heydarnoori**

Department of Computer Engineering, Sharif University of Technology, Tehran, Iran  
[heydarnoori@sharif.edu](mailto:heydarnoori@sharif.edu)

**Pooyan Jamshidi**

Department of Computing, Imperial College London, London, UK  
[p.jamshidi@imperial.ac.uk](mailto:p.jamshidi@imperial.ac.uk)

Technical Report

Automated Software Engineering Group (<http://ase.ce.sharif.edu>)

Department of Computer Engineering

Sharif University of Technology

October 2015

## Citation

A. Balalaie, A. Heydarnoori, P. Jamshidi, Microservices Migration Patterns, Technical Report No. 1, TR-SUT-CE-ASE-2015-01, Automated Software Engineering Group, Sharif University of Technology, October, 2015 [Available Online at <http://ase.ce.sharif.edu/pubs/techreports/TR-SUT-CE-ASE-2015-01-Microservices.pdf>]

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Pattern Meta-model and Specification Template</b>	<b>2</b>
2.1	Pattern Specification Template . . . . .	3
<b>3</b>	<b>Migration and Re-architecture Patterns</b>	<b>3</b>
3.1	MP1-Enable the Continuous Integration . . . . .	3
3.2	MP2-Recover the Current Architecture . . . . .	4
3.3	MP3-Decompose the Monolith . . . . .	6
3.4	MP4-Decompose the Monolith Based on Data Ownership . . . . .	7
3.5	MP5-Change Code Dependency to Service Call . . . . .	8
3.6	MP6-Introduce Service Discovery . . . . .	9
3.7	MP7-Introduce Service Discovery Client . . . . .	10
3.8	MP8-Introduce Internal Load Balancer . . . . .	10
3.9	MP9-Introduce External Load Balancer . . . . .	11
3.10	MP10-Introduce Circuit Breaker . . . . .	12
3.11	MP11-Introduce Configuration Server . . . . .	13
3.12	MP12-Introduce Edge Server . . . . .	14
3.13	MP13-Containerize the Services . . . . .	15
3.14	MP14-Deploy into a Cluster and Orchestrate Containers . . . . .	16
3.15	MP15-Monitor the System and Provide Feedback . . . . .	17
<b>4</b>	<b>Selecting and Composing Migration Patterns</b>	<b>18</b>
<b>5</b>	<b>Adding New Patterns to the Repository</b>	<b>18</b>

# 1 Introduction

Microservices is a cloud-native architecture through which a software system can be realized as a set of small services. Each of these services are capable of being deployed independently on a different platform and ran in their own process while communicating through lightweight mechanisms like RESTFull APIs. In this setting, each service is a business capability which can use various programming languages and data stores [1].

Migrating current on-premise software systems to microservices introduces many benefits including, but not limited to, different management of availability and scalability for different parts of the system, ability to utilize different technologies and avoid technology lock-in, reduced time-to-market, and better comprehensibility of the code base [2]. Furthermore, the migrated system can make use of the elasticity and better pricing model of the cloud environment, and therefore, providing a better user-experience for its end-users [2]. Although microservices presents many benefits, it introduces distribution complexities to the system for which new supporting components, e.g., Service Discovery, are needed. The difficulties around the decomposition of a system into small services, and monitoring and managing these services are amongst the other important factors that make migrating to microservices a non-trivial and cumbersome task.

Migrating to cloud, and in particular migrating through cloud-native architectures, like microservices, is a multi-dimensional problem and a non-trivial task [3]. Hence, in the absence of a well-thought methodology, it can be changed to a trial-and-error endeavor which can not only waste a lot of time, but also lead to a wrong solution. Furthermore, regarding the variation of factors like the requirements, the current situation, the skills of team members, etc., in different companies and scenarios, a unique and rigid methodology would not be enough. Thus, instead of a one-fit-all methodology, a Situational Method Engineering (SME) [4] approach is needed.

The first step towards the SME approach is to populate a method base or pattern repository with reusable process patterns or method chunks which we designate them as migration patterns in this report. It is important to note that each of these migration patterns should conform to a predefined meta-model. To this end, using our previous experience in SME [5] and defining migration patterns [6], in this report, we generalize our experience in migration to microservices [2] and similar practices in the state-of-the-art of microservices as migration patterns. These patterns are merely focused on *Migration Planning* phase. We enriched each step in the migration process with the precise definition of the corresponding situation, the problem to be solved, and the possible challenges of the proposed solution. After that, we transferred them to a pattern template which its parts has one-to-one correspondence with the meta-model elements. Parts of these patterns describes exactly why we need supporting components, e.g., Service Discovery, in our architecture, and what is the requirements for their introduction. Additionally, we provide some solutions and advices for decomposition of a monolithic system to a set of services, and defining the current and the target architectures of the system as a roadmap for migration planning. Still, we provide some clues about the containerization of services and their deployment in a cluster. Keeping the system in a stable state after applying a pattern and performing one architectural change at a time were the most important factors in determining the granularity of patterns.

Having a suitable pattern repository, a method engineer can select required patterns using selection guidelines and construct a bespoke methodology by composing the selected patterns.

The rest of this report is organized as follows: In Section 2, we present the meta-model and the pattern specification template conforming to that meta-model which used for pattern's documentation. Section 3 elaborates the devised migration patterns in detail. Section 4 explains migration pattern selection and the procedure of creating a methodology by composing the selected patterns. Finally, the process of adding new patterns to the current pattern repository is discussed in Section 5.

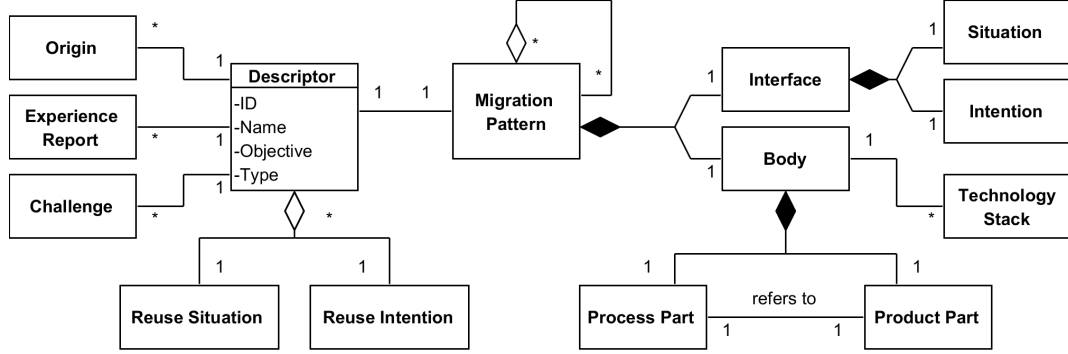


Figure 1: Meta-model

## 2 Pattern Meta-model and Specification Template

In order to be retrieved effectively from the repository and get reused, migration patterns should adhere to a meta-model. Additionally, the presence of a meta-model can help to expanding the repository by providing a mean for defining a new pattern in a standard manner. To this end, we reused the meta-model which has been proposed in [7] (see Figure 1). In this meta-model, each method chunk consists of the following parts:

- a *Descriptor* which is used for pattern selection and contains the name, ID, objective and type of the method chunk. Additionally, it provides references to the methodologies or experience reports from which the method chunk has been originated. Furthermore, it explains the situation wherein the method chunk can be reused and the intent behind its reuse.
- a *Body* which has a process part describing the solution that the method chunk is suggesting and a product part explaining the artifacts, if any, that should be created upon the execution of the process part.
- an *Interface* which provides information the situation in which the method chunk is suitable and its high-level objective.

In [8], a project development knowledge frame, namely the *reuse frame* has been proposed which is used for populating values for the *Reuse Situation*. The proposed reuse frame is way too general for our specific domain, i.e. migrating to microservices, and we find it not helpful in our particular case. Consequently, we proposed a set of architectural factors and a set of operational factors which can be used to specify the migration patterns' reuse situation. In Table 1, we briefly describe these factors. Additionally for *Reuse Intention* we chose to use the linguistic approach proposed in [9] which use a clause with a verb and a target for expressing the intention.

In order to make it more suitable for our migration patterns, we did some minor changes to the meta-model. First, we changed the name of method chunk class to migration pattern. Second, we enriched the descriptor by adding zero or more *Challenges* to it. As each migration pattern causes the system to undergo some changes and each change can introduce some side effects or challenges, we decided to put these challenges as a determining factor in pattern selection. Third, some of the migration patterns are related to introducing new supporting components or need some tools for their realization which are already have implementations in real world. As a result, we added the *Technology Stack* class to the body of the migration pattern.

Table 1: Migration pattern selection factors

Factor	Description
Architectural Factors	
Scalability	Increasing Scalability of an application by scaling out its services
High Availability	Increasing Availability of an application by replicating its services
Fault Tolerance	Decreasing the chance of failure in an application and providing means for handling failures effectively
Modifiability	Increasing the ability to change an application with the least side effects and without affecting its end-users
Polyglot-ness	Enabling an application to use different programming languages and data stores
Decomposition	Re-architecting an application to a set of services
Understanding	Perceiving the current situation of an application
Visioning	Deciding on the final situation of an application after migration
Operational Factors	
Dynamicity	Enabling an application to change in runtime without affecting its end-users
Resource-efficiency	Decreasing the amount of resources which are needed for an application's deployment
Deployment	Facilitating an application's deployment process and removing deployment anomalies
Monitoring	Enabling an application to be monitored in runtime effectively

## 2.1 Pattern Specification Template

Each migration pattern is documented in a pattern template. The patterns' title is a combination of their ID and Name joining with a "-". The patterns' type can be either atomic or aggregate. As can be seen in Section 3, most of the part's names in the pattern template use exactly the names in the meta-model. Nonetheless, there are some exceptions which are as follows:

- The word *Context* is used instead of *Situation* in the pattern's interface.
- The word *Intention* is replaced with *Problem* since it reflects the intention of the pattern in a question form.
- The *Solution* part is a replacement for the *Process Part* and *Product Part* collectively.
- The *References* part is a combination of Origins and Experience Reports.
- The *objective* part is removed as we find it overlapping with the *Reuse Intention*

## 3 Migration and Re-architecture Patterns

### 3.1 MP1-Enable the Continuous Integration

**Type** atomic

**Reuse Intention** Build the Continuous Integration pipeline

**Reuse Situation** Deployment

**Context**

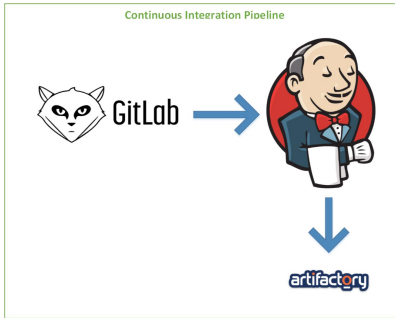


Figure 2: MP1-Enable the Continuous Integration

There is a working software system, and the team responsible for it has decided to migrate this system to microservices.

#### **Problem**

Considering that by adopting microservices, the number of services will be increased, how can we always have available production-ready artifacts? how to prepare the system for introducing the Continuous Delivery?

#### **Solution**

The first step towards the Continuous Delivery [?] is to set up the Continuous Integration [?]. Continuous Integration automates the build and test process and helps to always have production-ready artifacts. Normally, a Continuous Integration pipeline contains a code repository, an artifact repository and a Continuous Integration server. First, each service should be placed in a separate repository which helps to have a more clear history and separates the build life cycle of each service. Then, a Continuous Integration job should be created for each service. Each time a service's code repository changes, the job should be triggered. The job's responsibility includes fetching the new code from the repository, running the tests against the new code, building the corresponding artifacts and pushing these artifacts to the artifact repository. Failure in doing each of these steps should terminate the job from proceeding and informs the corresponding development team of the occurred errors. This team should not do anything else until addressing the reported errors. One simple rule in Continuous Integration is that new changes should break the system's stability and should pass all of the predefined tests.

#### **Technology Stack**

Gitlab, Artifactory, Nexus, Jenkins, GoCD, Travis, Bamboo, Teamcity

#### **References**

### **3.2 MP2-Recover the Current Architecture**

**Type** atomic

**Reuse Intention** Create Migration Plan initial state

**Reuse Situation** Understanding

#### **Context**

There is a working software system, and the team responsible for it has decided to migrate this system to microservices. In order to plan the migration, the team need to know the current system architecture which either does not exists or is rather obsolete.

#### **Problem**

What is the big picture of the system? What high-level information is sufficient for planning the migration to microservices?

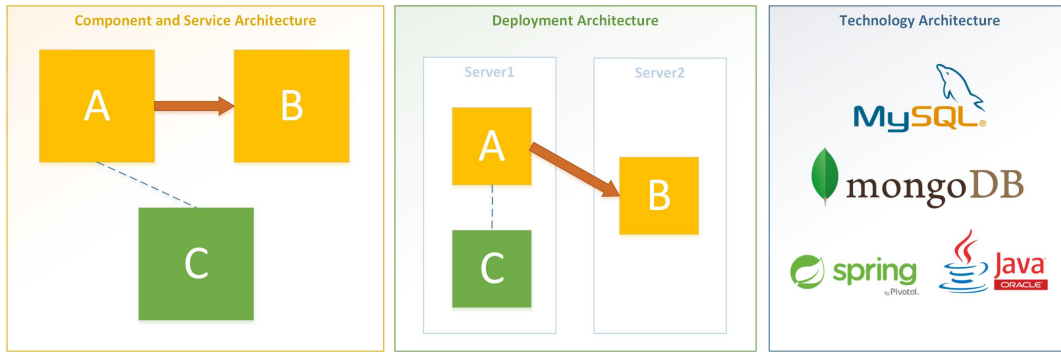


Figure 3: MP2-Recover the Current Architecture

### Solution

Talking about software architecture, people tend to imagine a bunch of formal diagrams; however, in practice, a helpful software architecture is a set of important things in the system communicated through some textual or visual artifacts. It is a good practice to keep these artifacts as simple as possible, thereby everyone can understand them easily using some basic hints. The following items are important in planning for migration to microservices:

- *Component and Service Architecture*: Using these two types of architecture together is not accidental. In practice, we find it more useful to illustrate these architectures in one informal visual artifact so everyone could easily grasp the inner structure of the system by a quick glance. Indicating the service calls direction is important as it could clearly separate service providers and service consumers. Furthermore, it could provide some clues about the dynamics of the system. Try to understand the inner domain of each component by considering their entities and their overall business logic. Discuss these items with the developers responsible for that specific component to identify additional details. Perceiving the domain of the system is important since it is required for decomposition of the system to small services.
- *Technology Architecture*: Understanding the current technology stack is important as it could help the team to identify existing libraries that could facilitate the migration. These new libraries include both the ones which can be used as supporting components in microservices that work well with the current technology stack, i.e. a service discovery for Java, and ones that facilitate non-architectural migration aspects, i.e. data migration tools. Enumerate any programming languages, database technologies, middleware technologies and third-party libraries that has been used in the project.
- *Deployment Architecture and Procedure*: microservices comes in a close relationship with the Continuous Delivery that is a radical change in software deployment. Understanding the current deployment architecture and procedure will help the team to gradually move towards the Continuous Delivery.

Do not document every bit of details as they are going to be changed in a near future, therefore, it is a waste of time. Try to involve developers and operations teams in process of understanding the system architecture since it is more time-efficient and they are needed for the rest of the migration. In this way, a common understanding of the system will be established. This information will be consolidated in the team members minds and communicated orally that is far better than an unread architecture document. Furthermore, a good atmosphere for collaborating between developers and operations can be established that is a good start

for a DevOps spirit. Put these artifacts in a place so that everyone in the team could see them. Be open to the comments since they may reveal some important details about the system.

## References

### 3.3 MP3-Decompose the Monolith

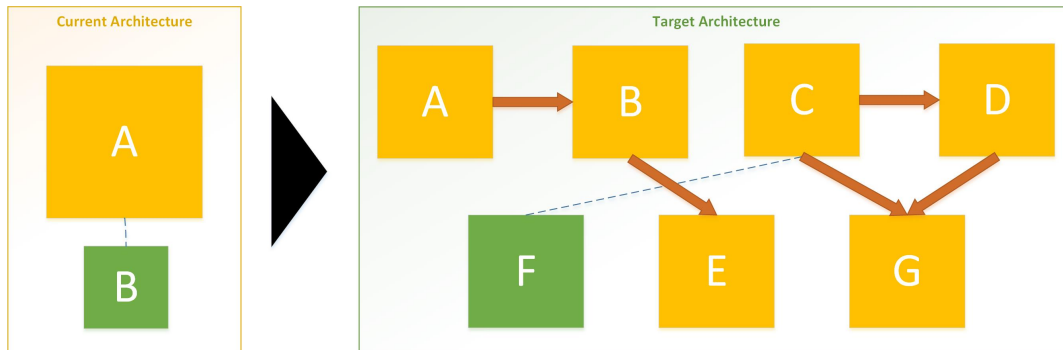


Figure 4: MP3-Decompose the Monolith

**Type** atomic

**Reuse Intention** Re-architect a System to a set of services using Domain-driven Design

**Reuse Situation** Polyglot-ness, Decomposition, Modifiability

**Context**

There is a monolithic software system with a complex domain which has at least one of the following attributes:

- Due to the complexity of the system, comprehensibility of the source code is low.
- Different parts of the system have different non-functional requirements, e.g. scalability.
- Every change in the system need a whole redeployment, and all parts of the system are not equal in terms of change frequency.

**Problem**

How to decompose the system into smaller chunks? How big these chunks should be?

**Solution**

Domain-Driven Design (DDD) should be used to identify sub-domains of the business that the system is operating in. Then, each sub-domain can constitute a Bounded Context (BC) that is a deployable unit. The one-to-one correspondence between sub-domains and BCs happens in greenfield projects. In contrast, due to the existing limitations in legacy systems, it is not always possible. Nevertheless, since the candidate system is going to benefit from microservices, it is strongly suggested that the responsible team plan to migrate the system like it is a greenfield project. It is true that this type of planning can introduce a lot of changes to the system, but it is the way that the migration could be most beneficial. Executing the plan incrementally is another concern, but at first, the team should have a good destination so that after migration the business could see the advantages.

DDD is a good option for initial decomposition of the system. Further decomposition can be happened as a result of either different change frequency rate or different non-functional requirements for different parts of a BC. Additionally, DDD can be applied on a sub-domain to decompose it to smaller chunks.



The size of the BCs is nothing that can be recommended or suggested. It totally depends on the systems requirements. At the beginning, they can be as large as the corresponding domain. As time goes on, requirements will change, and the previous boundaries or BC sizes may not be appropriate anymore. It is recommended to start with low number of services, e.g. two or three, and incrementally add more services as the system grows and the team understand microservices and the systems requirements better.

#### Challenges

A bad system decomposition can lead to performance penalties due to chatty services or unresolved challenges that lead the team to the process of migration.

#### References

### 3.4 MP4-Decompose the Monolith Based on Data Ownership

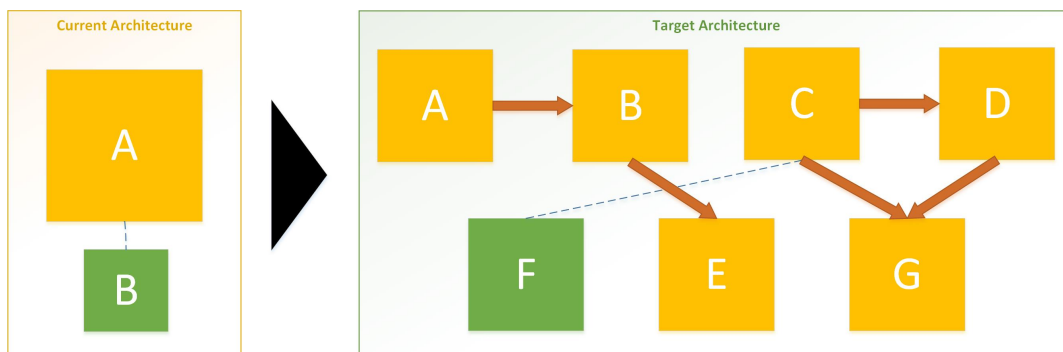


Figure 5: MP4-Decompose the Monolith Based on Data Ownership

**Type** atomic

**Reuse Intention** Re-architect a System to a set of services using Data Ownership

**Reuse Situation** Polyglot-ness, Decomposition, Modifiability

**Context**

There is a monolithic software system with not a complex domain which has at least one of the following attributes:

- Due to the complexity of the system, comprehensibility of the source code is low.
- Different parts of the system have different non-functional requirements, e.g. scalability.
- Every change in the system need a whole redeployment, and all parts of the system are not equal in terms of change frequency.

**Problem**

How to decompose the system into smaller chunks? How big these chunks should be?

**Solution**

Decompose the system based on the ownership of data. Find different cohesive sets of data entities that can be grouped together as a unit and can have a unique owner. Package each group and their corresponding business logic into a service. Each entity can be modified or created just through its owner that is its corresponding service. Other services can have a copy of an entity that they are not own, but they should be careful about its staleness and should synchronize their copy in an appropriate way. Further decomposition

can be happened as a result of either different change frequency rate or different non-functional requirements for different parts of a service.

The size of the services is nothing that can be recommended or suggested. It totally depends on the entities exist in the system. For a system with not a complex domain it would not be more than four or five services.

#### Challenges

This pattern is suitable when the domain of the system is not complex and the data entities can be grouped easily. In a large domain that has multiple sub-domains, applying this pattern could be confusing and time-consuming and even can lead to a not appropriate decomposition.

#### References

### 3.5 MP5-Change Code Dependency to Service Call




Figure 6: MP5-Change Code Dependency to Service Call

**Type** atomic

**Reuse Intention** Transform code-level dependency to service-level dependency

**Reuse Situation** Decomposition, Modifiability

**Context**

A software system has been decomposed to a set of small services to use microservices architectural style. There is a component in the system that is acting as a dependency to another services or components.

**Problem**

When it is appropriate to change this code-level dependency to service-level dependency? And when it is not?

**Solution**

Try to keep the services code as separate as possible. When services shares code as a dependency there is a high chance that a service developer can fail the build process of another service by changing the shared code. In some cases, like the code that is shared as a common library, e.g. string manipulation library, which rarely changes, it is reasonable to share code. However, sharing internal entities or interface schemas should be prohibited since changing them would force another dependent services to change immediately even though they may want to change gradually. In cases that sharing is not a good idea, it is a good practice to share that piece of functionality as a service. This service could be either a completely isolate service or a part of one of the dependent services. Another reason that could result in changing a shared library to a service could be different scalability needs between the shared library and the dependent services. In this case, by separating them in different services, they could be scaled independent of each other.

### Challenges

Sometimes separating libraries from their dependents and using a service call instead of a method call could introduce performance issues. Although performance issues can be handled using careful caching mechanisms, it adds another layer of complexity to the dependent service.

### References

## 3.6 MP6-Introduce Service Discovery




Figure 7: MP6-Introduce Service Discovery

### Type atomic

**Reuse Intention** Introduce Dynamic Location of services' instances using Service Discovery

**Reuse Situation** Scalability, High Availability, Dynamicity, Deployment

### Context

A software system has been decomposed to a set of small services to use microservices architectural style, and each of these services has one or more instances deployed in the production environment. The number of instances can be changed dynamically and each of them can be deployed in different systems.

### Problem

How services can locate each other dynamically? How an Edge Server or a Load Balancer knows the list of instances of a service to which it can route the traffic?

### Solution

Setup a Service Discovery which stores service instances addresses. Each service registers itself during initiation. The removal of an instance from the registry can be triggered through either not receiving a periodic heartbeat from the instance or by the instance itself during termination. Having a list of available services instances, an Edge Server, a Load Balancer or another services can locate their desired services dynamically through Service Discovery.

Service Discovery when introduced, is a vital component of the system since the communication between different parts of the system depends on the availability of the Service Discovery. Thus, replication strategies should be leveraged as a high availability mechanism.

### Challenges

The rest of the system are coupled to this component for communication among each other. Thus, Service Discovery, if not correctly implemented, could become a single point of failure.

### Technology Stack

Eureka, Consul, Apache Zookeeper, etcd

## References

### 3.7 MP7-Introduce Service Discovery Client




Figure 8: MP7-Introduce Service Discovery Client

**Type** atomic

**Reuse Intention** Facilitate Dynamic Location of services' instances using Service Discovery Client

**Reuse Situation** Scalability, High Availability, Dynamicity, Deployment

**Context**

A software system has been decomposed to a set of small services to use microservices architectural style, and a Service Discovery has been set up. The number of instances of each service can be changed dynamically and each of them can be deployed in different systems.

**Problem**

How the Service Discovery knows a new instance has been deployed? How it can know an instance has been terminated?

**Solution**

Each service should know the address the Service Discovery and registers itself during initiation. Then, a periodic heartbeat should be sent from each instance towards the registry so as to keep the instance in the available instances list. Instance termination can be done through either not sending heartbeat anymore or explicitly informing the registry of the instance termination.

**Challenges**

The drawback is that the client should be implemented for all of the programming languages in use and its bad implementation can complicate the service code.

**Technology Stack**

Eureka is a Service Discovery which has a Java client implementation for its server version.

**References**

### 3.8 MP8-Introduce Internal Load Balancer

**Type** atomic

**Reuse Intention** Introduce Load Balancing between instances of a service using Internal Load Balancer

**Reuse Situation** Scalability, High Availability, Dynamicity

**Context**

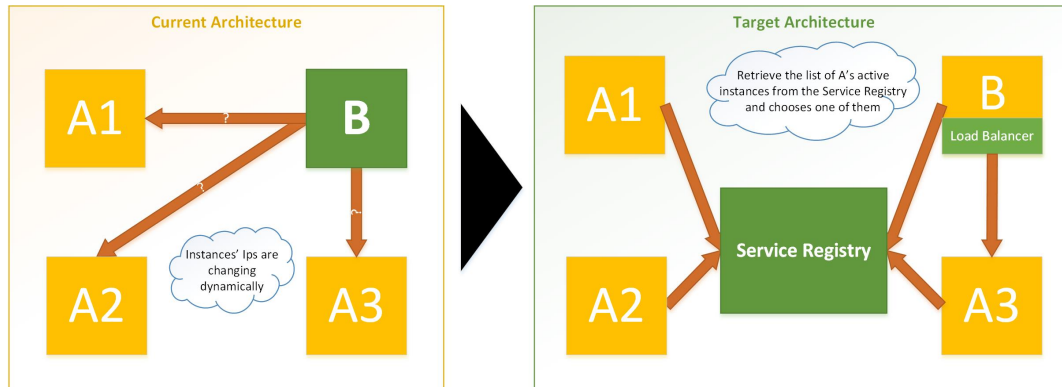


Figure 9: MP8-Introduce Internal Load Balancer

A software system has been decomposed to a set of small services to use microservices architectural style, and a Service Discovery has been set up. In the production environment, numerous instances of each service exist. Each service can be a client of the rest of the services in the system.

#### Problem

How to balance the load on a service between its instances based on the conditions of the client? How to load balance a service without setting up an external load balancer?

#### Solution

Each service, as a client, should have an internal load balancer which fetches the list of available instances of a desired service from the Service Discovery. Then, this internal load balancer can balance the load between the available instances using local metrics, e.g. response-time of the instances.

An internal load balancer removes the burden of setting up an external load balancer and brings in the possibility of having different load balancing mechanism in different clients of a service.

#### Challenges

The downside is the need for creating an internal load balancer for different programming languages in use which can be integrated with the Service Discovery. Additionally, the load balancing mechanism is not centralized.

#### Technology Stack

Ribbon is an internal load balancer for Java that works well with Eureka, a Service Discovery.

#### References

### 3.9 MP9-Introduce External Load Balancer

**Type** atomic

**Reuse Intention** Introduce Load Balancing between instances of a service using External Load Balancer

**Reuse Situation** Scalability, High Availability, Dynamicity

#### Context

A software system has been decomposed to a set of small services to use microservices architectural style, and a Service Discovery has been set up. In the production environment, numerous instances of each service exist. Each service can be a client of the rest of the services in the system.

#### Problem

How to balance the load on a service between its instances with the least changes in the service's code? How to have a centralized load balancing approach for all of the services?

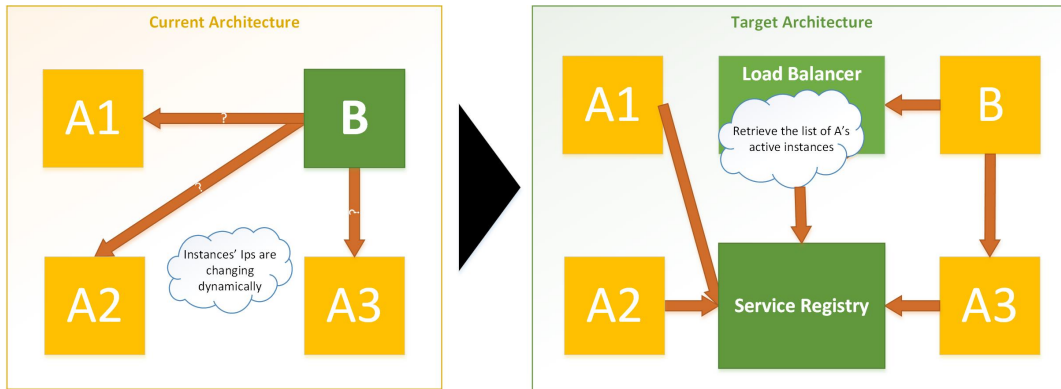


Figure 10: MP9-Introduce External Load Balancer

### Solution

An external load balancer should be set up which retrieves the list of available instances from the Service Discovery and uses a centralized algorithm for balancing the load between instances of a service. This load balancer can act either as a proxy (not recommended) or an instance address locator. Another solution is to choose a Service Discovery that has a built-in support for load balancing and can act as a load balanced address locator.

### Challenges

The downside is that local metrics, e.g. response-time of the instances, cannot be used to improve the load balancing. Furthermore, different clients cannot have different load balancing strategy. Additionally, when the load balancer acts as a proxy a high available load balancing cluster is needed.

### Technology Stack

Amazon ELB, Nginx, HAProxy, Eureka

### References

## 3.10 MP10-Introduce Circuit Breaker

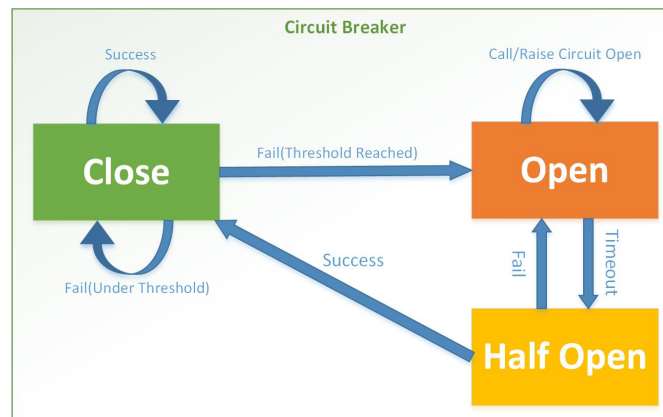


Figure 11: MP10-Introduce Circuit Breaker

**Type** atomic

**Reuse Intention** Introduce Fault Tolerance in inter-service communication using Circuit Breaker

**Reuse Situation** Fault Tolerance, High Availability

**Context**

A software system has been decomposed to a set of small services to use microservices architectural style. Some of the end-user requests need inter-service communication in the internal system.

**Problem**

How to fail fast and not wait until reaching a service call time-out when calling a recently unavailable service?  
How to make the system more resilient in the case of calling an unavailable service by providing reasonable response on behalf of the service?

**Solution**

The service consumer should use a Circuit Breaker [10] when calling the service provider. When the service provider is available, this component would not do anything (Close Circuit state). It monitors the recent responses from the service provider and will act appropriately when the number of failure responses passes a predefined threshold (Open Circuit state). The corresponding action could be returning a meaningful response code or exception, returning latest cached data from the service provider if it is reasonable and so on to the caller. After a specific timeout, in order to check the service provider availability, the component will try to access the service provider again (Half-open Circuit state) and in case of a successful attempt the state will be changed to Close Circuit. Otherwise, the state will be modified to Open Circuit.

**Challenges**

Recognizing the appropriate response in the Open Circuit state could be a bit challenging and should be coordinated with the business stakeholders. Furthermore, if the response is not just an exception, the service provider team should certify the possibility of returning that response on their behalf.

**Technology Stack**

Hystrix

**References**

### 3.11 MP11-Introduce Configuration Server

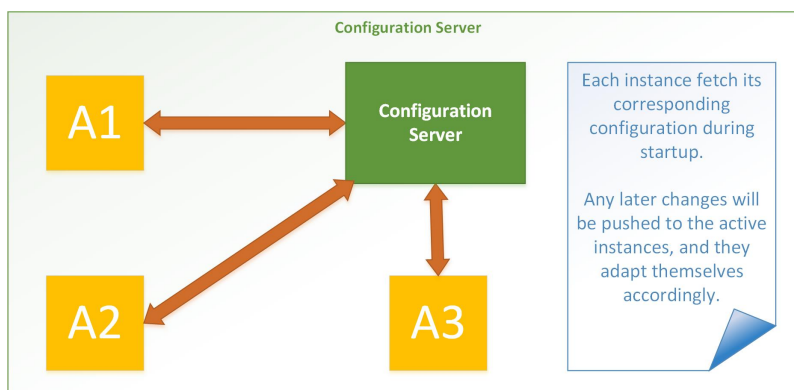


Figure 12: MP11-Introduce Configuration Server

**Type** atomic

**Reuse Intention** Change a system's configuration at runtime using Configuration Server

**Reuse Situation** Modifiability, Dynamicity, Deployment

### Context

A software system has been decomposed to a set of small services to use microservices architectural style. Each service has numerous instances running in the production. The list of available instances is available through a Service Discovery.

### Problem

How to modify the running instances configuration without redeploying them?

### Solution

There should have two separate repositories for storing source codes and software configurations. Although there may be a need for synchronizing these repositories when some changes happen in the configuration keys, they should evolve independently of each other. Furthermore, any change to the configuration repository should be propagated to the corresponding running instances and they should adapt themselves accordingly.

### Challenges

The configuration propagation endpoints in the services and adapting strategy should be implemented for all of the programming languages in use.

### Technology Stack

Spring Config Server, Archaius

### References

## 3.12 MP12-Introduce Edge Server

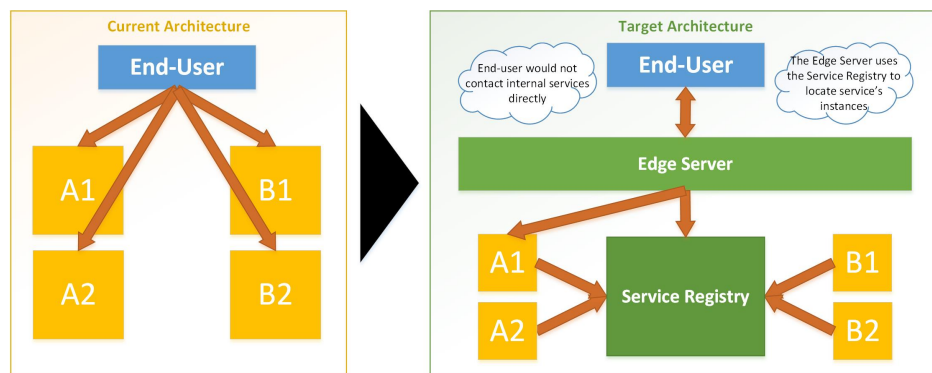


Figure 13: MP12-Introduce Edge Server

### Type atomic

**Reuse Intention** Enable Dynamic Re-routing of external requests to internal services using Edge Server

**Reuse Situation** Modifiability, Dynamicity

### Context

A software system has been decomposed to a set of small services to use microservices architectural style. Due to ease of service initiation in the system, new services can be introduced easily, and existing services can be re-architected based on new requirements.

### Problem

How to hide the internal service structure complexity and evolution from the end-users? How to monitor the overall status and usage of a service in production?

### Solution

There should be another layer of indirection in the system as a front door of the system. This layer does dynamic routing based on a predefined configuration. The service instances addresses for routing the incoming



traffic can be either fixed and hard-coded or fetched from a Service Discovery. End-users will depend on this layers interface, and thus, the internal structure changes would not affect them and will be handled through new routing rules. Furthermore, since all of the traffic will passes through this layer, it is the best place to monitor overall usage of the services.

#### Challenges

As all the traffic passes through this layer, and in order to remove the single point of failure, this layer should be replicated through load balancing mechanisms.

#### Technology Stack

Zuul

#### References

### 3.13 MP13-Containerize the Services

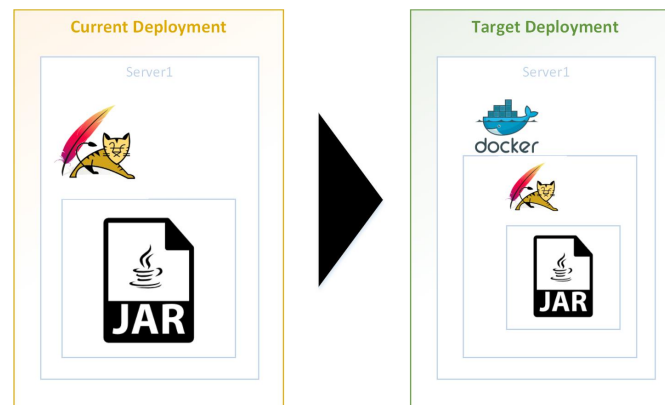


Figure 14: MP13-Containerize the Services

**Type** atomic

**Reuse Intention** Have the same behavior in development and production environment using Containerization

**Reuse Situation** Resource-efficiency, Deployment

#### Context

A software system has been decomposed to a set of small services to use microservices architectural style, and a Continuous Integration pipeline is in place and is working. Each service needs a specific environment so as to run correctly which is set up manually or through a configuration management tool. The differences between the development and production environments cause some problems like same code producing different behaviors in these two environments. All to all, the deployment of so many services in production has become either a cumbersome or complex task.

#### Problem

How to make the development and production produce the same results for the same code? How to remove the complexity of configuration management tools or difficulties of manual deployment?

#### Solution

As each service may need different environment for deployment, a solution could be deploying each service in a virtual machine in isolation and provide it with its desired environment using configuration management tools. The downside is that due to Virtualization, a lot of resources are wasting for service isolation and configuration management is another layer of complexity in deployment. Compared to the Virtualization,

the Containerization is more lightweight and it can remove the need for configuration management tools since there are a lot of ready images in the central repositories containing different applications, and any further configuration can be done in new images building stage.

In order to make it happen, add another step to the Continuous Integration pipeline to build container images and store the images in a private image repository. After that, these images can be ran in both production and development environments producing the same behavior. Each service should have its container image creations configuration and the script for running any other required services containers inside its code repository.

It is a good practice to add environment variables as a high priority source for populating the software configuration. In this way, the configuration keys that can have different values in different environments, e.g. database URL or any credentials, and they can be injected easily in container creation phase. Having a list of required environment variables for running a service in its code repository is a good practice and can make anyone aware of these changing variables.

### Challenges

Containerization could introduce computational overhead in comparison to deployment directly to an OS in favor of a lightweight, isolated and reproducible environment. Furthermore, the development environment should be adapted to embrace containers.

### Technology Stack

Docker

### References

## 3.14 MP14-Deploy into a Cluster and Orchestrate Containers

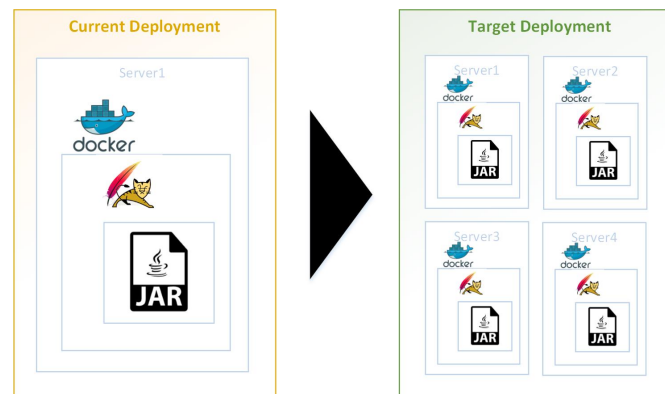


Figure 15: MP14-Cluster Management and Container Orchestration

### Type atomic

**Reuse Intention** Deploy service instances' container images in a cluster using cluster management tools

**Reuse Situation** Resource-efficiency, Deployment

### Context

A software system has been architected based on microservices architectural style, and a Continuous Integration pipeline is in place and is working. Therefore, a production-ready container image is available for the deployment of each service. Furthermore, there exists a large number of services and thereby their deployment and re-deployment is complex, cumbersome and unmanageable.

### Problem

How to deploy a service's instances into a cluster? How to orchestrate the deployment and re-deployment of all of the services with the least effort?

**Solution**

There should exist a system that can manage a cluster of computing nodes. This management system should be able to deploy the services container images on-demand, with specified number of instances and on different nodes. Additionally, it should handle the failure of instances and restart the failed nodes or instances in case. Still, it should provide a mean for auto-scaling of the services. Having an internal name resolution strategy could be a good feature as some services such as Service Discovery should be identified using name instead of an IP address.

In order to effectively orchestrate the deployment process, the cluster management tool should provide a mean to define the deployment architecture of services declaratively. Having a declarative deployment architecture, any additional effort for deployment, e.g. auto-scaling and failure-management of the deployed services, should be delegated to the cluster management tool.

**Challenges**

None

**Technology Stack**

Mesos+Marathon, Kubernetes

**References**

### 3.15 MP15-Monitor the System and Provide Feedback

**Type** atomic

**Reuse Intention** Monitor the running services' instances and Provide feedback to the development team

**Reuse Situation** Monitoring, Modifiability

**Context**

A software system has been architected based on microservices architectural style. The whole system is being ran on a cluster of containers with each service having a number of instances in production.

**Problem**

How to monitor the underlying infrastructure? How to use the gathered data to re-architect the system by providing feedbacks to the development team?

**Solution**

Each service should have its independent monitoring facility owned by the Ops part of the services team. These facilities includes the required components to gather the monitoring information, e.g. CPU and RAM usage, and sending them to a monitoring server. Therefore, the following components should be added to each service containers image and be configured during either creation of container images or creation of actual containers. Then, in the monitoring server, these information should be parsed and aggregated into a structured information and stored somewhere that can be queried efficiently, e.g. an indexing server. Having a pool of timely monitoring information, a visualization tool can be used to get an overall status of the system. This information helps the Dev part of the service's team to refactor the architecture in order to remove the performance bottlenecks or other detected anomalies.

**Challenges**

None

**Technology Stack**

Collectd+Logstash+ElasticSearch+Kibana

**References**



they can be added to the current factors.

## References

- [1] M. Fowler and J. Lewis, “Microservices.” <http://martinfowler.com/articles/microservices.html>, March 2014. [Last accessed 3-October-2015].
- [2] A. Balalaie, A. Heydarnoori, and P. Jamshidi, “Migrating to cloud-native architectures using microservices: An experience report,” in *1st International Workshop on Cloud Adoption and Migration (Cloud-Way*, (Taormina, Italy), September 2015.
- [3] P. Jamshidi, A. Ahmad, and C. Pahl, “Cloud migration research: A systematic review,” *IEEE Transactions on Cloud Computing*, vol. 1, pp. 142–157, July 2013.
- [4] B. Henderson-Sellers, J. Ralyt, P. Agerfalk, and M. Rossi, *Situational Method Engineering*. Springer-Verlag Berlin Heidelberg, 2014.
- [5] M. Gholami, M. Sharifi, and P. Jamshidi, “Enhancing the open process framework with service-oriented method fragments,” *Software & Systems Modeling*, vol. 13, no. 1, pp. 361–390, 2014.
- [6] P. Jamshidi, C. Pahl, S. Chinenyeze, and X. Liu, “Cloud migration patterns: A multi-cloud architectural perspective,” in *10th International Workshop on Engineering Service-Oriented Applications (WESOA)*, 2014.
- [7] B. Henderson-Sellers, C. Gonzalez-Perez, and J. Ralyte, “Comparison of method chunks and method fragments for situational method engineering,” in *19th Australian Conference on Software Engineering (ASWEC)*, pp. 479–488, March 2008.
- [8] I. Mirbel and J. Ralyt, “Situational method engineering: combining assembly-based and roadmap-driven approaches,” *Requirements Engineering*, vol. 11, no. 1, pp. 58–78, 2006.
- [9] N. Prat, “Goal formalisation and classification for requirements engineering,” in *Requirements Engineering: Foundation for Software Quality*, (Spain), p. 1, 1997.
- [10] M. Nygard, *Release It!: Design and Deploy Production-Ready Software*. Pragmatic Bookshelf, 2007.